

FinPersona-Bench: Where the Novelty Stands and How to Position Each Contribution

What the revised paper is

The revised FinPersona-Bench is a mechanistic study of mandate salience with behavioural validation. It asks three questions about an LLM agent that is given a standing behavioural mandate and left to trade over a long rollout: does the mandate lose control of the agent's decisions; what in the model carries that loss; and what does re-grounding the mandate do about it.

Thesis. Re-grounding is a directive amplifier whose sign depends on the fit between the directive's content and the market regime, and its effect is located in the model rather than inferred from behaviour alone. The instrument is built so that it stands even if the thesis returns a null.

How behaviour is measured. Objectively, step by step, against a mathematically defined fundamental value that the agent never sees, in a synthetic market with calm, crash, bubble and sustained-bull regimes that is audited for leakage and validated against the stylized facts of asset returns.

How adherence is scored. With mandate-conditional regret and literature-anchored target bands, normalised against a trivial-policy envelope, plus a target-free measure of mandate salience.

How cause is isolated. A causal arm design separates what a re-injected mandate does through its content, through its imperative force and through its position in context.

How mechanism is established. A programme on open-weight models (attention to the mandate span, a linear adherence direction, causal ablation and add-back, gated attribution) explains the behavioural result and drives an adaptive re-grounding controller.

Unique contributions as they will stand

- **The instrument.** A synthetic market in which every step of a single LLM agent is scored against a fundamental value that is generated by a calibrated fundamentalist-chartist process, never displayed, audited for leakage with thresholds fixed in advance, phase-controlled through randomised regime schedules, validated against the stylized facts of asset returns, and independent of the agent's own trades; with a sustained-bull control that separates adherence to the mandate from being right about value.
- **The metrics.** Mandate-conditional regret (an oracle conditioned on the mandate, so a compliant conservative agent incurs zero regret), targets as literature-anchored bands, a stated-target versus unstated-target axis, a trivial-policy envelope every number is normalised against, and a target-free mandate-salience measure.
- **The causal decomposition.** Declarative placebo, directive placebo, wrapper-only, mandate, swapped mandate and no-mandate arms, with mandate-wording variants, which separate content from imperative force from position on a ground-truth behavioural readout, and yield the content sign reversal as their consequence.
- **The mechanism.** Behavioural coverage on open-weight models that gates the internals work; attention to the mandate span with attention-sink control; a linear mandate-adherence direction validated against the behavioural metric; causal ablation and add-back read on the behavioural metrics with no token moved; gated logit attribution for the reversal; and an adaptive re-grounding controller triggered by the internal adherence signal, compared against always-on re-grounding, the injection-frequency sweep and a score-based trigger.

- **The persona measurement.** MBTI and Big Five personas sampled across the trait space, with the persona-text share and the directive share of allocation behaviour measured rather than asserted, so the framing rests on behavioural consistency, not psychometrics.
- **The statistics.** Model-level inference (mixed effects with random effects for model, persona and seed; BH correction across the metric family), so every headline is a claim about models, not about runs.

The table below takes each contribution in turn: what the literature already has, what is still ours, and how to word and design the claim so that it is defensibly novel. Papers are linked where they are mentioned.

1. Contribution by contribution

Contribution as the revised paper will state it	What the literature already has	What is still ours	How to position it to build novelty
Mandate Salience Decay as the object of study: a standing behavioural instruction losing control of behaviour over a long rollout.	The concept is owned several times over. Meta's Proactive Memory Agent defines 'behavioral state decay' (information still in context that no longer governs behaviour). Li et al. measure instruction drift across dialogue turns. Dongre gives a dynamical model of drift with a restoring force and an intervention term. When Attention Closes measures a goal-accessibility ratio falling 27 to 48 percent over 50 turns. Arike et al. measure goal drift in an LLM trading agent. All of these have real accumulating context.	Not the concept and not the word 'formalise'. What is ours is the measurement: every step is scored against a mathematically defined hidden fundamental value and a persona target band, with a resolvability rule and a mandate-conditional oracle, in the same environment for every model. Every prior paper scores drift with judges, classifiers, questionnaires, self-reports or task accuracy.	Do not say 'we formalise MSD'. Say: 'we measure mandate adherence objectively, per step, against a mathematically defined ground truth that no other drift study has' and cite Meta , Li , Dongre and Arike as the concept's owners in the first paragraph. Design that carries it: hidden V with published leakage audit, band targets, mandate-conditional regret, trivial-policy envelope. Null to reject: drift is bounded noise (Dongre); we test it with windowed statistics against a constant-exposure null on the stateful arm. If the stateful grid shows decay, the claim upgrades to 'the first ground-truth-scored demonstration of decay in a behavioural allocation task'.
Behaviour is distinct from reasoning: agents violate the mandate while their stated reasoning and their restatement of the mandate stay coherent (knows-but-violates).	Owned. Kruthof's DriftBench : violation rates 8 to 99 percent with 97 percent recall of the violated constraint, and an explicit argument against the 'forgetting' reading. CLQT : stated analysis and allocation disagree (coherence gap +0.30 backtest, +0.23 live). Song et al. : persona prompts move questionnaire answers, not generation. Kocielnik et al. : persona prompting stabilises self-reports without aligning behaviour. Everitt et al. separate goal-directedness from capability.	The financial allocation instance, measured on the same step: forked restatement probe (what the agent says its mandate is) against the target-share action it takes, scored against ground truth. With the stateful agent, both sides of the dissociation (accessibility falling, behaviour holding or not) are measured in one rollout, which no paper in the corpus does (When Attention Closes has accessibility, Kruthof has recall, neither has the behavioural ground truth).	Position it as 'the accessibility-behaviour dissociation measured in allocation space over a real rollout', not as a new concept. Use Kruthof's knows-but-violates as the hypothesis being tested, cite it by name, and report the restatement-probe accuracy beside the behavioural metric in the same table. Under the stateless arm say plainly that every violation is knows-but-violates by construction; the stateful arm is where the claim lives. Design that carries it: forked restatement probe, stateful agent with the mandate in the system prompt and a rolling context, M1 attention to the mandate span.
A synthetic market that decouples observable price from a fundamental value the agent does not see, giving objective, falsifiable ground truth.	Partly owned. Hashimoto et al. : LLM traders in a limit-order-book market with a hidden GBM fundamental. Machine Spirits : 15 LLMs in Hommes-style markets with a known fundamental (disclosed to the agents; price endogenous). ABIDES : exogenous oracle fundamental, no LLM agents. Stylized-facts validation (Hashimoto , TwinMarket) and leakage probes (Machine Spirits) exist. Fundamental-chartist mispricing is Franke and Westerhoff's ; the stylized facts are Cont's .	A leak-audited, phase-controlled, stylized-facts-validated single-agent environment with per-step, agent-independent ground truth: mispricing generated by a calibrated fundamental-chartist process, randomised regime schedules so phase and time are not confounded, a sustained-bull control that separates 'adhered to the mandate' from 'was right about value', published L1 to L5 leakage thresholds on the rendered observation, a resolvability rule, trivial-policy and oracle normalisation, and thresholds fixed in advance. Order-book simulators cannot give a single LLM agent a per-step scoring truth that is independent of that agent's own trades.	Claim it as a measurement instrument, never as market realism. The sentence: 'the first environment in which every step of a single LLM agent is scored against a hidden value that is (i) generated by a calibrated behavioural-finance process, (ii) audited for leakage with published thresholds, and (iii) agent-independent'. Cite Hashimoto and Machine Spirits as the existence proofs and ABIDES as the oracle idea, then state the three things they do not have (audit, single-agent per-step truth, resolvability). With the audit table beside it, 'hidden' can stand; say 'not directly displayed' only where an audit number requires it. Put the checklist and the leakage table in the main paper, not an appendix: they are the contribution.
Regimes and metrics: calm,	Partly owned. CLQT has a mandate-alignment term and a	The specific combination: targets as literature-	Lead with mandate-conditional regret as the metric

Contribution as the revised paper will state it	What the literature already has	What is still ours	How to position it to build novelty
crash (three severities), bull trap with a real top, sustained bull; band-MAS, drawdown, mandate-conditional regret.	Discipline axis on real data with costs. Arike's action score measures deviation of budget share from the goal (it counts saving the budget as drift, which mirrors the conservative persona). Regime-switching market design is standard. Drawdown is drawdown. Jiang, Peng and Yan show traits explain about 5 percent of human equity share, so targets are stipulated bands, not derived points.	anchored bands rather than points, mandate-conditional regret (distance from the best action the mandate allows, so a conservative agent holding cash in an undervalued market incurs zero regret), a stated-target vs unstated-target axis (Track A vs Track B), crash severities that move drawdown, topped vs un-topped bubble strata, and a trivial-policy envelope every number is normalised against. No paper has the persona-fair regret or the stated/unstated axis.	contribution: 'a regret that is fair across mandates because the oracle is conditioned on the mandate', and show in one figure that RG-style absolute scoring ranks a perfectly compliant conservative agent as irrational. Present Track A vs Track B as a designed contrast that turns the tautology objection ('the target mirrors the persona') into a measured quantity. Drop 'Caricature Index' as a name; call it drawdown. Cite CLOT's Discipline axis and Arike's action score as the nearest relatives and state the difference in one line each.
Breadth: a broad roster of frontier and open-weight models, every persona, every regime, replicated seeds and decodes.	Owned as a practice. ContextEcho 23 models; PERSIST 25 to 29; DriftBench 7; Menon , SEQUOR , Machine Spirits , Multi-IF and AI-Trader all report strong model dependence of the same kind.	Not model dependence itself. Ours is the combination of breadth with a fixed, ground-truth-scored environment and model-level statistics (mixed effects with random slopes by model, BH across the metric family), plus a $t = 0$ separability gate with criteria fixed in advance whose per-family pass rate is itself a result.	Do not list 'the effect is model dependent' as a contribution. Use the breadth as the statistical design: 'effective N is models, not runs', and make the gate pass-rate table a headline ('which model families can even instantiate a persona mandate at $t = 0$ '). This turns a weakness the literature already knows into an instrument result nobody has reported for allocation behaviour.
Temporal structure: does the static-vs-re-grounded gap grow, settle, or shrink over the rollout, including a long-horizon control.	Owned and contradicted. Dongre : multi-turn drift settles at a bounded, noise-limited equilibrium with a restoring force. Laban et al. : degradation present from the second turn, dominated by variance. Multi-IF : the per-turn decline is front-loaded. SEQUOR : re-statement resets compliance and decay is not explained by context length. No paper supports a compounding law.	A clean test: with randomised regime onsets (phase and day are not confounded) and a stateful agent, the paper can be the first to test bounded-versus-compounding drift in a behavioural allocation task against a constant-exposure null, on the main horizon and on a long-horizon control.	Do not claim compounding. Position the temporal analysis as a test, stated in advance, of Dongre's bounded-equilibrium prediction: windowed statistics, stateful arm, phase controlled for. Report whichever way it goes; a confirmation of bounded drift in a new domain with ground-truth scoring is citable, a refutation is a headline. Cite Dongre , Laban , Multi-IF and SEQUOR in the hypothesis statement so the test reads as designed, not salvaged.
Re-grounding is not universally beneficial: the same re-injected mandate helps capital-preservation agents and hurts growth agents in an identical market (content sign reversal; the surviving headline).	The broad form is owned. Mehri et al. re-inject the user goal every turn and find model- and dimension-specific reversals. Dongre's reminders raise one model's drift equilibrium. The Always-On survey documents 'baseline beats memory' across studies. Kruhthof's checkpointing only partly reduces violation. PushBench's checklist reminder gives zero benefit. The experimental form is also owned: ContextEcho pairs an imperative anchor at the recency position with a length-matched control, ablates its content, and finds a 'compliance amplifier' that lifts even no-drift targets. Ross and Lo show stateless allocations follow the stated risk label, so opposed targets moving in opposite directions is the expected outcome.	The effect as a behavioural allocation against opposed, literature-anchored targets, in a ground-truth-scored market, over a long rollout, across a broad model roster, with the content isolated by the swapped-mandate arm (same wrapper, other persona's directive) and the directive force isolated by the directive placebo and wrapper-only arms. No paper has the content-vs-force-vs-position decomposition on a behavioural readout.	Position it as 'a directive amplifier whose sign is set by the fit between directive content and market regime', and make ContextEcho's compliance amplifier the explicit null: if the directive placebo and wrapper-only arms move allocation as much as the mandate, the effect is force, not content; if the swapped mandate moves allocation to the injected side, it is content. Report the four-arm plus swapped-mandate decomposition as the result, the sign reversal as its consequence. Cite Mehri , Dongre , ContextEcho and Ross and Lo in the setup. This is the one behavioural headline the literature leaves open; the decomposition is what makes it ours rather than a domain replication.
Placebo and control package: content versus text position versus directive force.	Owned as a design. ContextEcho : length-matched filler arm in every cell (41,921 evaluations), user-turn vs system-prompt placement, content ablation, 23 models. When Attention Closes Appendix H: three-arm design (analogous, not identical: prepended goal, fact recall, forced sliding window). Arike : token-distance ablation. Arditi et al. : random-suffix control.	The completeness and the readout: declarative placebo, directive placebo, wrapper-only, mandate, swapped mandate and no-mandate trader, plus mandate-wording variants, in every cell, on a ground-truth-scored behavioural task. ContextEcho cannot say whether its anchor works by content or by imperative force; our package answers exactly that question.	Never call the placebo 'a placebo control' as if it were new. Say: 'we extend ContextEcho's length-matched design with a directive placebo and a swapped-mandate arm so that content, imperative force and position are separately identified on a behavioural metric'. Cite ContextEcho , When Attention Closes and Arditi in the design paragraph. Design that carries it: all six arms in every cell, with declarative and imperative placebos as separate arms so that the two are never confounded.
Persona framework independence: MBTI, Big Five on the full grid with continuous trait sampling, and a numerical-	Owned and partly inverted. Kocielnik : Big Five self-reports are reliable but do not predict within-model behaviour. Song : persona effects are carried by lexical cues. PERSIST : trait instability. Ross and Lo : the stated	The quantification: how much of allocation is carried by a stated target (Track A, O3) versus persona vocabulary (Track B), measured on the same models, same environment, with the $t = 0$	Do not claim framework independence. Claim: 'we measure the persona-text share and the directive share of allocation behaviour and show the directive dominates, in line with Ross and Lo for LLMs and with Jiang, Peng and Yan for humans'.

Contribution as the revised paper will state it	What the literature already has	What is still ours	How to position it to build novelty
only (target stated) arm.	risk label carries 57 to 88 percent of predictive weight in LLM allocations. Jiang, Peng and Yan : in humans, Big Five traits explain about 5 percent of equity-share variance, so targets cannot be derived from traits.	gate and the target-free salience measure (persona share vs directive share of the surrogate). The finding that persona text is a weak allocation signal and the directive is a strong one is consistent with the human finance evidence and is measured, not asserted.	Present the O3 numerical-only result as the ceiling of what a stated target achieves and the Track B result as what persona vocabulary achieves. Rename 'psychometric' to 'behavioural consistency' throughout. Cite Kocielnik and Song in the setup as the reason the measurement is needed.
Re-grounding schedule: an injection-frequency sweep scored on the behavioural metric, and an adaptive controller that re-injects only when an internal adherence signal falls below threshold.	The sweep is owned: Li et al. swept injection probability for system-prompt repetition with an MMLU capability-cost axis and proposed split-softmax as a mitigation. Selective triggers are owned: Meta built a selective memory agent; CLQT ships a drift-warning trigger when its consistency score falls below 0.7.	The controller keyed to an internal drift signal (the M2 mandate direction) rather than to a behavioural score, compared on the same ground-truth metric against always-on, the k-sweep and a CLQT-style score trigger, with the token cost reported. A closed-loop method with behavioural validation is a method contribution; the sweep alone is a replication.	Present the k-sweep as the baseline curve, cite Li for it, and put the cost axis (context tokens) on the same plot as Li does. Present the controller as the contribution: 'the first re-grounding policy triggered by a mechanistic adherence signal and validated on a ground-truth behavioural metric'. Cite Meta and CLQT for selective triggering and state what differs: the trigger reads internals, not outputs, and the benefit is measured per regime, including the regimes where always-on re-grounding hurts.
Mechanistic programme: behavioural coverage on open-weight models (M0) gating attention to the mandate span with attention-sink control (M1), a linear mandate-adherence direction (M2), causal ablation and add-back with the prompt unchanged (M3), logit attribution on the mandate-to-SELL pathway (M4), and the rationale-linguistic analysis as behavioural validation (M5).	The methods are owned. Li : attention to the system prompt decays across turns. When Attention Closes : goal accessibility falls while recall holds; patching null at n = 30. Arditi : behaviour direction by difference-in-means, ablation and add-back. Chen et al. : persona vectors monitored during generation. Xiao et al. : attention sinks. Meng et al. : activation patching and logit attribution.	Tying these methods to a long-horizon, ground-truth-scored behavioural metric: attention and direction projection per step across the rollout, and a causal intervention whose effect is read on band-MAS and drawdown rather than on a probe accuracy. Nobody in the corpus connects the mechanistic signal to a behavioural allocation outcome in a real rollout, and nobody has the accessibility-vs-behaviour dissociation with both sides instrumented.	Position the M-series as 'mechanistic validation of a behavioural benchmark', not as new interpretability. Cite Li , Arditi , Chen , Xiao and Meng as the methods in one sentence each and then state the novelty: the readout is a ground-truth behavioural metric over a long rollout. Under the stateful arm the question is decay; under the stateless arm the same methods answer the content-vs-force question (what in the residual stream carries the directive), which is still unclaimed. Make M3 the centrepiece: an effect produced with no token moved is the cleanest answer to the 'content and position are confounded' critique.
Selective, event-triggered re-grounding as the prescription; transfer to medical triage and legal compliance.	Owned. Meta built selective injection; CLQT ships a score trigger and names mandate variation as future work. Menon et al. already run an ER triage environment (45 steps, 5 seeds).	Nothing as a prescription or a transfer claim. What is ours is the controller in the row above and the portable instrument design (hidden ground truth, mandate-conditional regret, leakage audit) that other domains could copy.	Delete the prescriptive and transfer sentences from the conclusion. Replace with one line: 'the instrument design, a hidden ground truth with a mandate-conditional oracle and a published leakage audit, is portable to any domain with a computable counterfactual', and cite Menon for triage so the reader knows it exists.
The benchmark as a released artefact: a long, paired, judge-free, ground-truth-scored behavioural instrument, released with its generator, audits, baselines, oracles, rendered prompts and provenance hashes.	ContextEcho is longer (9,643 turns) but has no ground-truth-scored task. CLQT is a closed-loop financial benchmark on real data with costs but no hidden truth. DriftBench and Multi-IF are judge- or rule-scored. No corpus paper publishes a leakage audit with thresholds fixed in advance on the observation the agent actually sees.	The longest paired, generating-function-scored, judge-free behavioural instrument in the corpus, with trivial-policy, constant-mix, momentum, mean-reversion, value-oracle, mandate-conditional-oracle and observables-oracle baselines per cell, leakage and phase-clock audits with thresholds, and provenance hashes in every row.	Claim it explicitly as an artefact contribution: 'an audited benchmark whose thresholds were fixed before the runs' is rare in this literature and reviewers reward it. Put the audit tables and the baseline envelope in the paper body or a named appendix, and release generator, audits, baselines, rendered prompts and hashes together.

2. The big picture: three pillars the paper can stand on

After the literature review the paper does not have a novel concept. It has a novel instrument, a novel causal decomposition, and a novel mechanism result. Written that way, the paper is still in the competition. Written as 'we formalise MSD; MSD compounds; re-grounding is the cure', it collides with eight papers that did each of those things first.

Pillar	What it consists of	Why it is ours
1. The instrument	The synthetic market: calibrated fundamentalist-chartist mispricing on a hidden value, randomised regimes, sustained-bull control, leakage and phase-clock audits with published thresholds, stylized-facts checklist, band targets, mandate-conditional regret, trivial-policy envelope, target-free salience measure.	No existing LLM-agent environment gives a single agent a per-step, agent-independent, audited ground truth. Hashimoto and Machine Spirits have fundamentals but no audit and no single-agent scoring truth; CLQT has costs and discipline but no hidden truth.
2. The causal decomposition	Declarative placebo, directive placebo, wrapper-only, mandate, swapped mandate and no-mandate arms; mandate-wording variants; stated-target vs unstated-target tracks; model-level statistics. Result: the content sign reversal and how much of it is content, force, or position.	ContextEcho , Li , Arike and When Attention Closes each have one or two of these arms; none has the full decomposition on a behavioural ground-truth metric, and none can say whether the amplifier is content or force.
3. The mechanism	Stateful agent with real accumulating context; forked restatement probe; M0 behavioural coverage on open-weight models; M1 attention with sink control; M2 adherence direction; M3 ablation/add-back read on the behavioural metrics; M4 gated attribution; M5 behavioural validation; adaptive controller triggered by M2.	Li , When Attention Closes , Menon and Kruthof each hold half of the accessibility-vs-behaviour result; no one has both sides instrumented in one long rollout with a ground-truth behavioural readout, and no one has a controller driven by an internal signal validated that way.

3. Wording: what to say and what to stop saying

Stop saying	Say instead	Reason
We formalise Mandate Salience Decay.	We measure mandate adherence objectively, per step, against a mathematically defined hidden value; decay is one of the hypotheses we test.	The concept is Meta's , Li's , Dongre's and Arike's . The measurement is ours.
MSD compounds over time (4.4x).	We test whether the gap grows, settles or shrinks, against Dongre's bounded-equilibrium prediction, with phase and time un-confounded.	Four papers predict bounded drift; the claim needs the temporal test, not an assertion.
A placebo control shows the effect is content, not position.	A six-arm decomposition separates content, imperative force and position; the swapped-mandate arm identifies content.	ContextEcho's compliance amplifier is the null; a single declarative placebo cannot reject it.
Re-grounding is a structural prerequisite; selective triggers are the prescription.	Re-grounding is a directive amplifier whose sign depends on content-regime fit; we ship and evaluate a controller triggered by an internal adherence signal.	Meta and CLQT already built selective triggers; a measured controller is a contribution, a prescription is not.
Hidden fundamental value.	A fundamental value that is not displayed and is audited for leakage (L1 to L5, thresholds published).	Reviewers will check; the audit table is the defence and the contribution.
Caricature Index; Rationality Gap.	Drawdown; mandate-conditional regret.	CI is drawdown; RG penalised compliant agents and was saturated by a trivial policy.
Psychometric stability; framework independence.	Behavioural consistency; the persona-text share versus the directive share of allocation.	Kocielnik , Song and PERSIST undercut the psychometric reading; the share is measurable.
Transfers to triage and legal.	The instrument design is portable to any domain with a computable counterfactual.	Menon already runs triage.

4. Papers the revision must cite

None of the following is cited in the current draft.

Paper	Why it must be cited	Link
Arike et al. 2025	Direct competitor: behavioural goal drift of LLM traders; imperative elicitation reduces drift; token distance is not the driver.	Arike et al. 2025: arxiv.org/abs/2505.02709
ContextEcho 2026	Closest design analogue: length-matched filler, imperative anchor at recency, content ablation, 23 models; the compliance-amplifier null.	ContextEcho 2026: arxiv.org/abs/2605.24279

Paper	Why it must be cited	Link
Li et al., COLM 2024	System-prompt repetition swept over injection probability with a cost axis; cross-turn attention decay; split-softmax. Basis for M1 and the k-sweep.	Li et al., COLM 2024: arxiv.org/abs/2402.10962
Kruthof, DriftBench 2026	Knows-but-violates with 97 percent recall; rejects the forgetting framing; the hypothesis our restatement probe tests.	Kruthof, DriftBench 2026: arxiv.org/abs/2604.28031
Dongre, Drift No More 2025	Bounded drift with a restoring force; reminders lower but do not remove the equilibrium; the temporal null.	Dongre, Drift No More 2025: arxiv.org/abs/2510.07777
When Attention Closes 2026	Goal accessibility decays while recall holds; prepended re-injection does not help; the other half of the dissociation.	When Attention Closes 2026: arxiv.org/abs/2605.12922
Meta, Proactive Memory Agent 2026	Defines behavioral state decay; selective versus always-on injection.	Meta, Proactive Memory Agent 2026: arxiv.org/abs/2607.08716
CLQT 2026	Closed-loop financial benchmark with a mandate-alignment axis, a drift-warning trigger, a stating-versus-doing gap; names mandate variation as future work.	CLQT 2026: arxiv.org/abs/2606.29771
Menon et al. 2026	Frontier models hold a system-prompt goal for 30 steps under pressure; user-turn directives override system goals for some families; ER triage already run.	Menon et al. 2026: arxiv.org/abs/2603.03258
Ross and Lo 2026	Stateless allocations follow the stated risk label; makes opposed-target movement the expected outcome and motivates Track A vs Track B.	Ross and Lo 2026: arxiv.org/abs/2604.23837
Kocielnik et al. 2026; Song et al. 2026; PERSIST 2025	Big Five does not predict LLM behaviour; persona effects are lexical cues; trait instability. The reason the persona-share measurement is needed.	Kocielnik et al. 2026: arxiv.org/abs/2606.12730 Song et al. 2026: arxiv.org/abs/2509.10078 PERSIST 2025: arxiv.org/abs/2508.04826
Mehri et al. 2026	Per-turn goal re-injection with model- and dimension-specific reversals.	Mehri et al. 2026 (TACL): arxiv.org/abs/2507.20152
Hashimoto et al. 2025; Machine Spirits 2026	LLM traders with a fundamental value: the existence proofs the instrument claim must name.	Hashimoto et al. 2025: arxiv.org/abs/2510.12189 Machine Spirits 2026: arxiv.org/abs/2604.18602
Jiang, Peng and Yan, JFE 2024	Traits explain about 5 percent of human equity share; targets are stipulated bands, not derived points.	Jiang, Peng and Yan, JFE 2024: doi.org/10.1016/j.jfineco.2024.103776
Arditi et al. 2024; Chen et al. 2025; Xiao et al. 2024; Meng et al. 2022	Direction extraction, ablation and add-back, persona vectors, attention sinks, activation patching: the M1 to M4 methods and controls.	Arditi et al. 2024: arxiv.org/abs/2406.11717 Chen et al., Persona Vectors 2025: arxiv.org/abs/2507.21509 Xiao et al., Attention Sinks 2024: arxiv.org/abs/2309.17453 Meng et al. 2022 (ROME): arxiv.org/abs/2202.05262
Laban et al. 2025; Multi-IF 2024; SEQUOR 2026; Always-On Agents survey 2026	Multi-turn degradation is front-loaded and variance-dominated; re-statement resets compliance; baseline-beats-memory literature.	Laban et al. 2025: arxiv.org/abs/2505.06120 Multi-IF 2024: arxiv.org/abs/2410.15553 SEQUOR 2026: arxiv.org/abs/2605.06353 Always-On Agents survey 2026: arxiv.org/abs/2606.30306
Franke and Westerhoff 2012; Cont 2001	The mispricing process and the stylized facts the environment is calibrated and validated against.	Franke and Westerhoff 2012 (JEDC): doi.org/10.1016/j.jedc.2012.01.005 Cont 2001, stylized facts: doi.org/10.1080/713665670