

FinPersona-Bench — Position Statement and Revision Plan Update

Where the paper stands after the Cycle 2 verification and the full 47-paper literature read; what it can uniquely claim now; what it will be able to claim after revision; and what must change in the current Revision Plan for Resubmission. 22 August 2026. Source documents: FinPersona-Bench_Cycle2_Verification_RedTeam_Report_Aug2026 (verification, red-team, literature), FinPersona-Bench_Revision Plan for Resubmission (the plan being updated).

1. Where the paper stands today

The draft reviewed is docs/COLM_FinPersona_Bench_Current_Paper.pdf — the May post-rebuttal version of the April COLM submission (29 pp, App. A–L; byte-identical to COLM_FinPersona_Bench_Post_Rebuttal.pdf; no newer draft exists in the repository). Its four stated contributions (p. 2) are: (1) MSD formalised as a behavioural failure distinct from reasoning errors; (2) a benchmark decoupling observable price from hidden fundamental value; (3) across 18 LLMs, MSD "compounds over time" (4.4× crash gap) and re-grounding is "not universally beneficial"; (4) three diagnostics — placebo, Big Five/numerical-only ablation, injection-frequency ablation. The Conclusion (p. 10) promises "a mechanistic account of how mandate tokens lose influence as context accumulates" — the sentence the current revision plan is built around.

The numbers are sound; the story is not. Tables 1, 9, 10, 11 and Appendices E, F, G, L reproduce digit-for-digit from the raw runs. What fails verification is the mechanism narrative, three figures, the pooled crash headline, and the novelty positioning. One internal inconsistency in the draft points to the fix: §3.2 (p. 5) and Fig. 2 say the core mandate M is part of the persistent system message and that static agents get it "once at initialization", but App. B (p. 16) says M is "a concise, fixed behavioral directive used for memory re-grounding" — which is what the code does. Align §1/§3.2/Fig. 2 to App. B. The table below states, claim by claim, what the draft says and what the evidence supports.

Claim as written	Status	What the data support
Mandate Saliency Decay: mandates erode as market context accumulates	DOES NOT SURVIVE	Agents are stateless single-turn callers (Eq. 3 says so; the prose contradicts it). Nothing accumulates; the static arm never even contains the core mandate. What exists is a portfolio ratchet: per-step violation propensity is flat-to-declining while cash depletes. Prior art (Li 2024, Chen 2025, Xiao 2024, Kruthof 2026) independently predicts no attention-level decay in a fixed-length prompt.
Re-grounding restores a decayed mandate	DOES NOT SURVIVE as framed	The manipulation is mandate-ABSENT vs mandate-PRESENT on a stateless prompt. Nothing is re-injected. Honest description: directive present vs absent at the recency position.
4.4× compounding of the crash gap Q1→Q4	DOES NOT SURVIVE	Figure 3 is computed on days 1–100 only (binning bug, numerical_analysis.py:577–582) and relabelled. Corrected on the plotted statistic the ratio is ~9× but phase-locked to the scripted panic (local gap collapses in Q4). No corpus paper supports a compounding law.
Crash: static agents suffer 12.6% greater drawdown (p=0.002)	DOES NOT SURVIVE pooled; SURVIVES for ISFJ	Pooled: 12/18 models, every defensible model-level test p=0.10–0.24. ISFJ-scoped: 14/18, +7.5 pp, p=0.0038 — mechanism is the guardian staying out of the market (35–47% zero-trade runs), not panic-selling prevention.

Bidirectional sign reversal: 17/18 improve (ISFJ), 16/18 degrade (ENTJ)	SURVIVES as a fact; framing does not	Counts exact; model-level interaction $p=3.8e-5$; $dz \approx \pm 1.1$; replicates under OCEAN and $T=800$. But every mandate RAISES cash (+25.9/+9.6/+5.9 pp) — the behavioural effect is one-signed; the metric sign flips through opposed C_ideal targets. Content survives in magnitude ordering (protective >> growth) and in a null declarative placebo.
Placebo shows content, not position, drives the effect	WEAK	One model, 15 pairs; declarative vs imperative + wrapper confounded; placebo itself significant vs static ($p=0.041$). ContextEcho's 23-model finding that an imperative anchor lifts even no-drift targets ("compliance amplifier") is the unrejected null.
Hidden fundamental value as decoupled ground truth	PARTLY	V is recoverable to 0.17% from reported P/E ($EPS = V/15$). The generating-function scoring is sound; the word "hidden" is not.
Three psychometric metrics (MAS, CI, RG)	ONE SURVIVES WITH REPAIR	MAS_ISFJ $\equiv 1$ – cash; CI is maximum drawdown; RG is saturated by buy-one-share-then-HOLD (99.7) and its random baseline is 2/3 not 0.5. MAS and RG retain partial independence; CI and cash are one variable.
MBTI personas; Big Five replication as validation	DOES NOT SURVIVE as validation	Personas are not behaviourally separable at $t=0$; no persona approaches its target; Kocielnik 2026 shows Big Five is non-predictive of LLM behaviour; Song 2026 shows persona effects are lexical-cue-driven — which FinPersona's own O3 result fits.
LCR threshold predicts mini-model reversal; classifier P(intended persona)	DOES NOT SURVIVE	LCR has no cross-model validity ($r=0.13$, $p=0.60$); classifier is trained on both arms (mandate-echo detector).
Injection-frequency ablation (App. H) as a contribution	ANTICIPATED	Li et al. COLM 2024 swept injection probability with a capability-cost axis; App. H is a replication on a different axis.

1.1 Unique contributions we actually hold right now

1. A content-modulated directive effect at 18-model scale, exactly reproducible: appending a behavioural directive shifts allocation toward cash with magnitude ordered by content (protective >> growth >> neutral), a length-matched declarative placebo moves nothing (one model), and a numeric-target-only directive removes the aggressive-persona harm in all three models tested. Scored against opposed targets this yields the 17/18 vs 16/18 split (interaction $p=3.8e-5$). In the 47-paper corpus the injection-and-control LOGIC is owned by ContextEcho and Li et al.; the DOMAIN and READOUT (a behavioural allocation in a ground-truth-scored market, opposed targets, identical market) are not owned by anyone.
2. A persona-scoped stress result: capital-preservation directive reduces drawdown in 14/18 models in the crash ($p=0.0038$), with an honest mechanism (market avoidance).
3. The instrument: an 18-model $\times$ 3-mandate $\times$ 3-regime $\times$ 5-seed paired grid at 200 and 800 steps, scored against a generating function without an LLM judge, at a few cents per run, with every table reproducible from the released CSVs. Longest ground-truth-scored paired instrument in the corpus (ContextEcho is longer but has no generating-function task).
4. A vocabulary gradient (MBTI $\rightarrow$ OCEAN $\rightarrow$ numeric-only O3) that already shows trait language, not the target, carries the harm — a finding, once reframed honestly.
5. Two analyses produced by the review that are more interesting than the claims they replace: the portfolio-ratchet decomposition (per-step propensity vs integrated allocation) and the metric-refraction result (one cash variable, three author-placed targets). Owning these pre-empts the reviewer who would otherwise find them.

What is no longer a contribution and must be withdrawn: MSD-as-decay and the title; the 4.4x; the pooled crash/panic-selling failure mode; CI as a caricature measure; RG as defined; the psychometric framing; the placebo as a clean null; every classifier-derived number; the LCR mechanism; App. H's operational k recommendations.

2. Where the paper will stand after revision — two paths

There are two ways to revise, and they lead to different papers. Path A keeps the current single-turn architecture and reframes the paper around the effect the data already contain. Path B rebuilds the agent and environment so that the titular phenomenon — mandate salience decay under genuinely accumulating context — is instantiated and measurable, and runs the mechanistic programme as the plan designed it. Path B is larger, longer and riskier; it is also the paper the title promises and the one the COLM reviewers (VG7b, Gxm9) were asking for. Nothing from the existing grid is wasted under Path B: the current runs become the stateless control arm. Both paths share the same first step (the truth-reconciliation corrections, Section 3) and the same literature positioning.

2.1 Path B — rebuild so that MSD is actually instantiated (recommended if the budget allows)

What must be built for the MSD claim to be testable, and what the literature forces you to control for:

Requirement	Why	Design note
Accumulating context in the agent	Today each step is an independent call; the mandate is at the same token distance on day 200 as on day 1. Decay needs a transcript: at step t the model must see (a compressed record of) steps $1..t-1$, so the mandate genuinely sits further back as t grows.	200 steps $\times$ $\sim 1-2k$ tokens = 200–400k tokens — beyond most windows. Choose explicitly: rolling window (then you are replicating When Attention Closes' forced-closure regime in finance — cite and differentiate), summarisation memory (then the memory architecture is an axis — Meta, Always-On survey), or a shorter-step/longer-context design. The choice determines what "context accumulates" means and must be stated.
Mandate present in the static arm	Today the static arm never contains the core mandate, so nothing can decay and nothing is re-grounded.	Put M in the system prompt at $t=0$ for ALL arms; the memory arm additionally re-injects it. Then "re-grounding" is literally true.
Stateless control arm	To attribute any effect to accumulating context rather than to injection position/content.	The current grid IS this arm — keep it. Static-stateful vs static-stateless isolates context; memory-stateful vs memory-stateless isolates re-grounding under context.
Time decoupled from market phase	In the crash the scripted phases coincide with quartiles, so "compounds over time" is confounded with "responds to panic"; Q4 local gaps collapse in scripted stabilisation.	Randomise phase onsets across seeds, run long flat horizons (decay with no phase structure), or hold phase fixed while varying elapsed context. The rework of the synthetic environment is the place to do this.
Forgetting vs knows-but-violates	Kruthof 2026: models violate constraints they can restate (97% recall); Arike/Dongre/Menon show behavioural drift without context-loss. A decay claim must show the mandate is less ACCESSIBLE, not only less obeyed.	Forked restatement probe at intervals ($\sim 5\%$ extra calls); M1 attention accessibility (now measurable); M2 projection trajectory. Pre-register: decay = accessibility AND behaviour fall together; KBV = behaviour falls, accessibility does not.
Amplifier vs restoration	ContextEcho: an imperative anchor lifts even no-drift targets to ceiling. A re-grounding benefit must exceed what the same directive does with no decay present.	The stateless memory arm (already run) is exactly this baseline: benefit_stateful – benefit_stateless is the decay-specific component.
Content vs position vs imperative force	The placebo confounds declarative vs imperative and omits the wrapper; reviewers will ask.	Four-arm control (declarative, directive placebo, wrapper-only, mandate) + swapped mandate, on a subset — cheap, and the same arms Path A needs.
Environment repairs carried into the rebuild	P/E discloses V to 0.17%; RG saturated by buy-and-hold; C_ideal uncommunicated; δ -sweep near-vacuous; SMA/Table 2 mismatches.	Decouple EPS from V (or retract "hidden"); mandate-conditional regret with trivial-policy floors; communicated-target track; make δ actually move the price path; align prompt and Table 2.

Model-level inference	Effective N is models; seeds are reused.	Mixed-effects with random slopes by model; add 30–70B open models (also the M-series anchors).
-----------------------	--	--

Mechanistic programme under Path B: M1–M4 as the plan wrote them, now meaningful — M1 attention to the mandate span across the (real) rollout, static vs memory, sink-controlled (Li 2024; Xiao 2024), plus the position/content comparison; M2 mandate direction tracked per step (Arditi; Chen), validated against the action token/cash and the restatement probe (which now has variance); M3 ablate-early/add-back-late AND the all-layer mediation test (Arghal's single-layer null; Li's split-softmax lever); M4 gated attribution with causal tracing + gradient×activation (ROME supplies the former only). What the prior art predicts you will find: attention to the mandate declines across turns (Li — likely to replicate); behavioural decay for frontier models over 30–50 turns is weak or absent (Menon; When Attention Closes) and drift, where present, is bounded (Dongre) — so the interesting, publishable outcome is the DISSOCIATION (accessibility falls, behaviour holds or not, by model and regime), which the M-series is exactly built to show. Budget for that outcome; do not promise monotone compounding.

Claim if Path B succeeds (accessibility and behaviour both decay, re-grounding restores beyond the stateless amplifier baseline, M3 mediation holds): "Mandate Salience Decay is real, mechanistically grounded, and content-dependent" — the paper as titled, now supported, with a synthetic market that separates time from phase. Venue: COLM/ICLR main track. Claim if behaviour does not decay while accessibility does: "LLM financial agents keep obeying mandates they attend to less — drift is knows-but-violates, not forgetting, and reminders act as amplifiers" — a strong negative result against the field's decay narrative, main-track-worthy if clean. Cost: the rebuild (agent memory design, environment rework, re-running the grid statefully on 18+ models) is months, not weeks, plus the GPU M-series; and the title/abstract cannot be finalised until the stateful grid is in.

2.2 Path A — reframe on the current architecture (cheap, fast, smaller)

If the rebuild is not affordable this cycle, the paper reframes around the effect already measured. The revision then has one fork, decided by ~200 cheap simulations (the swapped-mandate + directive-placebo + wrapper-only arms). Everything else is the same in both branches.

	Branch A — content tracks (swap arm: ISFJ text into ENTJ agent moves cash toward 1.0; directive placebo < 30% of mandate effect)	Branch B — directive force dominates (effects follow imperative form, not content)
Headline claim	"Directive content, not position or form, steers stateless financial agents: a ground-truth-scored benchmark across 18 models." Content attribution is established; the opposed-target split becomes a demonstration of how scoring geometry can mislead.	"Any imperative at the recency position shifts stateless financial agents toward inaction; content modulates magnitude only." An honest measurement/negative-control paper about re-grounding as deployed practice.
Unique value	First content-attributed steering result in a behavioural allocation readout with generating-function ground truth; the compliance-amplifier null (ContextEcho) rejected in a new domain; communicated-vs-uncommunicated-target axis (new, unowned).	Largest-scale evidence that re-grounding is a generic caution shift refracted by metric targets; a cautionary design lesson for every mandate-adherence benchmark (CLQT, Arike, Meta); the same instrument.
Venue	COLM main track defensible (with model-level statistics throughout and the 16-paper positioning).	Findings track / strong agents-evaluation workshop (NeurIPS/ICLR workshop, ACL Findings tier).
Mechanistic account (both branches; M-series, rescoped — Section 3.2)	Attention to the injected directive span (position and content, not time), a mandate/risk direction whose projection shifts under injection and tracks injected content, and a mediation test showing that ablating the direction abolishes the directive	Same tools; the account becomes "imperative form, not content, is what the model's computation responds to" — the direction tracks directive force. Still a mechanistic finding, still unowned.

	effect and adding it reproduces the effect without moving a token. No corpus paper has a mechanistic account of directive-content steering in a financial decision; this answers VG7b-W1 and Gxm9 in the way the plan intended.	
Secondary contributions (both branches)	ISFJ-scoped crash result; the 200/800-step paired instrument with trivial-policy baselines; the ratchet and refraction analyses; the O3 vocabulary gradient; optional behavioural-trigger controller benchmarked against CLQT's $C_t < 0.7$ and Li's SPR-p.	Same.
What the title says	Drop "Longitudinal Psychometric Stability" and "Mandate Salience Decay". Candidate: "When Reminding Backfires: Directive Content Steers Stateless LLM Trading Agents".	Same family; "Reminders Are Not Neutral: ...".

Either branch is publishable and either is better learned now than in review. Under Path A the mechanistic track is retained but re-aimed at the directive-content effect (Section 3.2), because on a stateless architecture there is no decay to instrument. Both paths require the truth-reconciliation rewrite first.

2.3 Recommendation

Decide Path B vs Path A on budget and calendar, not on ambition — Path B is the better paper if it can be finished. Either way: (1) do the P0 corrections now — they apply to both and remove the twenty-minute forensic risk; (2) run the cheap four-arm + swapped-mandate package now — it is a required control under Path B and the decisive experiment under Path A; (3) if choosing Path B, fix the context-accumulation design (window / summary / horizon) and the time-vs-phase decoupling in the environment BEFORE re-running the grid, and keep the current stateless grid as the control arm; (4) write the pre-registration for M1–M4 with the dissociation outcome as an explicit, publishable result.

3. Changes required to the Revision Plan for Resubmission

The plan's diagnosis (§0) is right — "proven" and "salience" are doing work the evidence cannot support — and its instinct to add a mechanistic layer is sound. The one thing the plan does not say, and must, is that M1–M4 as written presuppose accumulating context and a mandate in the static arm, neither of which the current implementation has. So the plan needs a PRECONDITION section: under Path B, build the stateful agent and the phase-decoupled environment first (Section 2.1) and then M1–M4 run as designed; under Path A, re-aim M1–M4 at the directive-content effect (Section 3.2). The item-by-item table below gives the Path A verdict in the "New status" column and the Path B delta where it differs. The plan also targets a superseded draft (4.1×, Table 7, N=210/14 models) and must be re-based on the current PDF (4.4×, Table 11, N=270/18 — actually 263 in the crash).

3.1 Item-by-item

Plan item (plan priority)	New status (Path A) — and Path B delta	Why, and what replaces it
§0 "Convert to a mechanistic paper with behavioral validation"	KEEP the ambition; add the PRECONDITION. Path B: keep as written once the stateful agent exists.	State explicitly that the mechanistic account of DECAY requires accumulating context and a mandate in the static arm — build both (Path B) or re-aim at directive-content steering (Path A). The confound objection (VG7b-W1) is answered twice under either path — cheaply on 18 models by the swapped-mandate/directive-placebo/wrapper-only arms, and causally on open models by M3 with no prompt tokens moving.
M0 — behavioural coverage on open models (HIGH)	KEEP (HIGH); widen. Same under Path B.	Still the gate for the M-series and also the route to model-level power. Add the 30–70B tier (Llama-3.3-70B, Qwen2.5-72B,

		Gemma-3-27B) with the full metric suite — the 4–9B models currently analysed are exactly where the behavioural effect is weakest or reversed (Qwen2.5-7B harmful, Gemma-3-4B the exception), so anchors must be chosen where the effect exists. Foreground the Qwen language/behaviour dissociation.
M1 — attention decay across the rollout (HIGH)	Path A: KEEP, RE-AIMED at POSITION and CONTENT (no rollout to decay over). Path B: KEEP AS WRITTEN — now measurable — plus the position/content comparison and a stateless null.	Path A: per call, attention mass from the generation position (first action token) to the persona/finance span, the injected directive span, the placebo/wrapper span and the observation span — sink tokens excluded, depth-matched control spans (Xiao 2024; the end-of-prompt placebo span is NOT a matched control for a start-of-prompt span). Does the same directive draw more attention in user than system position (Li; ContextEcho App. I)? does it predict the action and covary with phase rather than step index? Path B: the same measurement across a genuine rollout, static vs memory, with the stateless run as the flat-line control; Li predicts cross-turn decline — expect it to replicate.
M2 — mandate direction (HIGH)	KEEP, validate against BEHAVIOUR (both paths). Path B: the per-step projection trajectory becomes the decay signal and the controller trigger the plan wanted.	Difference-in-means between guardian- and commander-prompted activations (Arditi 2024; Chen 2025: response-token average for extraction, last-prompt-token projection for monitoring). Validate against the first action token / cash fraction — NOT against MAS (circular with the cash readout) and NOT against the classifier (mandate-echo detector); the restatement probe is usable under Path B (it has variance there) but not Path A ($\approx 100\%$). Test whether injection moves the projection toward the injected CONTENT (ISFJ text in an ENTJ agent) or persona identity — the mechanistic twin of the swapped-mandate arm. Expect a noisy per-step signal (Chen within-condition $r=0.25\text{--}0.81$; Arghal linear probes 39% vs MLP 70%).
M3 — causal intervention (HIGH)	Path A: KEEP, RE-DESIGNED as a MEDIATION test. Path B: KEEP ablate-early/add-back-late AS WRITTEN, plus the mediation test. Pre-registered either way.	Mediation test: (i) ablate the M2 direction at ALL layers and positions (Arditi Eq. 4; Arghal's 0/456 single-layer result rules out single-layer interventions) in the memory arm — does the directive's effect on the action token / cash vanish? (ii) add the direction in the static arm — does it reproduce the memory-arm shift with no tokens moved? (iii) Li's split-softmax boosting attention to the mandate span — same effect? Controls: random direction of matched norm, self-patch, matched pairs and seeds. Under Path B, ablating early and restoring late in a genuine trajectory additionally tests whether late restoration reverses accumulated drift. When Attention Closes App. D.5 ($n=30$, one model, fact recall) is a risk signal, not a settled null; commit to publishing a null.
M4 — mini-model reversal attribution (MEDIUM)	KEEP, GATED and re-tooled (both paths).	Gate unchanged: an open model must reproduce a reversal first (Qwen2.5-7B is currently the only open one; GPT-5-mini breaks the "mini" pattern; the pooled crash effect is n.s. at model level, so frame as explaining a model-specific reversal, not a class effect). Tools: causal tracing (ROME-style corrupt-and-restore on the directive span — ROME has NO direct logit attribution; fix the citation) plus gradient $\times$ activation attribution (Templeton-style) to the action-enum token logit; quantity q is not a one-token target.
M5 — reposition LCR as behavioural validation (LOW)	MODIFY: demote to a 2-model case study; retrain the classifier	LCR: precision 0.50, recall 0.33 across 18 models; GPT-5-mini and 5.4-mini silently excluded. Classifier must be retrained static-arm-only on held-out models or its numbers dropped. Neither can validate M2/M3; the behavioural readout (action token / cash) does.
Scale OCEAN to full grid (HIGH)	CUT as validation; keep existing data as the vocabulary gradient	Kocielnik 2026: Big Five is non-predictive of LLM behaviour, so OCEAN cannot validate; Song 2026: persona effects are lexical cues. The O1/O2/O3 gradient is already the useful object — report it as such.
Derive C_{ideal} from literature (HIGH)	MODIFY	Naive derivation backfires: Jiang/Peng/Yan give single-digit-pp effects ($R^2 \approx 0.05$) and MBTI omits Neuroticism — derived targets would collapse the spread and the effect. Do: report stipulated AND derived targets side by side; ADD a communicated-target

		track (state the target in the prompt, as O3 already does) vs the uncommunicated revealed-preference track — this turns the tautology into a designed contrast and is a genuinely new axis; publish the $t=0$ separability baseline.
Redesign RG as mandate-conditional regret (HIGH)	KEEP	Add trivial-policy floors (buy-and-hold 99.7, all-cash 79.8, random 66.7, constant-mix, "BUY iff $P/E < 15$ "); $EPS = V/15$ sensitivity {10,15,25}. If the P/E-rule saturates it, retire RG honestly.
Re-run k sweep on MAS (HIGH)	KEEP, demote to LOW	Li et al. 2024 already ran the p-sweep with a capability-cost axis; this is a replication on a different axis, not a contribution. Cite Li; remove the operational k recommendations.
Adaptive re-grounding controller (HIGH, depends on M2)	KEEP as P4; two triggers	Primary trigger behavioural ($ C_t - \text{target} $ over threshold) — works on all 18 models, no activations; secondary internal-loop variant using the M2 projection on open models (the plan's idea), reported as a comparison. Benchmark both against CLQT's $C_t < 0.7$ trigger and Li's SPR-p with a capability/cost axis (Li make it mandatory). Moderate risk of "no better than baselines".
Scale π_{placebo} to the full grid (MEDIUM)	PROMOTE to #1 and expand into P2	This is the decisive experiment, not a spot check: four arms (declarative placebo, DIRECTIVE placebo — imperative/delimiter/length-matched but behaviourally irrelevant, WRAPPER-ONLY — the "ACTIVE MEMORY REFRESH" framing with no mandate, mandate) plus the SWAPPED-MANDATE arm (ISFJ text into ENTJ agent and vice versa) and a wording dose-response. Pre-registered predictions. ~500–2,000 cheap sims. Informative either way.
Statistical overhaul (MEDIUM)	PROMOTE to P0	Plus: model-level tests, effect sizes, CIs, BH; withdraw the pooled crash claim or label it directional; rescope to ISFJ; $N=263$.
Port to ABIDES (MEDIUM)	DEFER; keep the retraction. Path B: fold into the environment rework only if multi-agent price formation is wanted.	Delete §3.1's "realism reintroduces subjectivity" — false (Hashimoto 2025, not ABIDES, is the LLM-agent existence proof). If the environment is being reworked anyway (Path B), the priorities are time-vs-phase decoupling, the P/E leak and a δ that actually moves the price path — before an order book.
Justify generator parameters; AR(1) sentiment (MEDIUM)	KEEP; under Path B becomes part of the environment rework (HIGH)	Also fix: SMA label (prompt says SMA60, paper SMA50), Table 2 (add portfolio state; remove MACD/dividend yield or add to prompt), the crash δ -sweep's near-vacuity (price-path MDD floor varies 0.6 pp across δ), and — for Path B — randomised phase onsets / long flat horizons so elapsed context and market phase are separable.
Fix the 4.1 $\times$ claim (LOW)	REPLACE: fix the Figure-3 binning bug and withdraw the multiplier	The current draft says 4.4 $\times$; the real defect is that the figure covers days 1–100 with axis labels claiming 1–200. Regenerate on the full horizon (the corrected figure is stronger: monotone ~9 $\times$ on the plotted statistic) but drop the multiplier and the "compounds" language; report phase-locking.
Fix Figure 4 / Table 7 (LOW)	REPLACE with the full correction list, HIGH	Current numbering: Fig 4's CRASH column is sign-inverted (prints worsens-counts; true ISFJ 14/18 — improves the paper); Fig 5's crash bars inverted; Tables 1/9/10 three sign conventions; $N=263$ and the dagger; placebo range 0.047; §4.4 self-reference; App. F cohort provenance (pilot, not main grid); App. H raw data absent from repo; Table 7 caption says pooled, numbers are flat-only; temperature 0.2 not 0.0; DeepSeek naming.
Rename "psychometric stability" (LOW)	PROMOTE to HIGH — retitle	Also rename the construct: "salience decay" is unsupported for a stateless agent (Kruthof: violation with 97% recall; Li; Chen). "Directive weighting" / knows-but-violates is what the corpus supports.
Scope the prescriptive claim (LOW)	KEEP	And remove "structural prerequisite"; note CLQT and Meta already ship selective triggers.
Promote $T=800$ (LOW)	KEEP	Verified exact from results_long_horizon; state the cohort.
Mandate wording sensitivity (LOW)	PROMOTE to HIGH — folded into P2	Paraphrase, abbreviation, delimiter on/off, imperative strength.
Figure 2 quartile footnote	SUPERSEDED	The footnote would document a statistic computed on the wrong

(LOW)		days; fix the bug instead.
-------	--	----------------------------

3.2 The mechanistic track under Path A (re-aimed) — summary

Under Path B the M-series runs as the plan wrote it (Section 2.1). Under Path A — the current stateless architecture — it is re-aimed from "how the mandate loses influence as context accumulates" (not measurable there) to "how an injected directive's content and position change the computation, and whether a mandate/risk direction mediates the behavioural effect" (measurable, unowned, and the direct answer to VG7b-W1/Gxm9).

Step	Question it answers	Method (corpus-grounded)	Pre-registered prediction	Reads out on
M0	Where does the behavioural content effect exist on open weights?	Full MAS/MDD/RG grid on 7B–72B open models; choose 2–3 anchors where the ISFJ/ENTJ effect is present	Effect present at $\geq 27B$; weak/reversed at $\leq 9B$	cash fraction, first action token
M1	Does the directive draw attention by POSITION (user vs system) and CONTENT, not by time?	Attention mass from generation position to directive / persona / placebo / observation spans; sink exclusion; depth-matched controls (Li 2024; Xiao 2024)	No step-index trend (null); user-position > system-position; attention-to-directive predicts action; covaries with phase	per-call attention shares
M2	Is there a mandate/risk direction, and does injection move the model along it in the direction of the injected CONTENT?	Difference-in-means guardian vs commander (Arditi; Chen); projection at last prompt token; swapped-mandate in activation space	Projection shifts under injection; tracks injected text, not persona identity (Branch A) or tracks imperative form (Branch B)	projection vs cash / action token
M3	Does the direction MEDIATE the behavioural effect?	All-layer directional ablation in memory arm; addition in static arm; split-softmax attention boost (Li); random-direction, self-patch, matched-pair controls; null committed to publication	Ablation abolishes $\geq 50\%$ of the directive effect; addition reproduces it with no tokens moved	cash / action token
M4 (gated)	Which part of the directive drives a model-specific reversal?	Causal tracing on the directive span + gradient $\times$ activation attribution to the action-enum logit, on an open model that reproduces a reversal	Reversal attributable to action-inciting tokens ("BUY/SELL"), not to the persona label	action logits
M5	Behavioural validation surface	Classifier retrained static-only, held-out models; LCR as 2-model case study	—	—

Sequencing: M0 runs with P3 (coverage). M1–M3 run after P0 (rewrite) and alongside P2, because M2/M3 ask mechanistically the same question P2 asks behaviourally (content vs identity vs directive force); reporting both is stronger than either. M4 only if its gate opens. Budget: one GPU-class machine, TransformerLens/nnsight on 8–27B (72B needs sharding); 4–6 weeks for M0–M3.

3.3 Items the plan does not contain and must add

- Truth reconciliation of the architecture prose (P0, 1–2 weeks, \$0): rewrite abstract/\$1/\$3.2/Fig 1–2/App. B to the stateless reality — the static arm contains no core mandate; nothing accumulates; portfolio state IS passed each step (Eq. 3/5 and Table 2 corrected). This is the single most consequential edit: a reviewer with the repo open finds it in twenty minutes.
- The 16-paper mandatory citation set with differentiation sentences (Cycle 2 report §8.7): Arike, Menon, Li (COLM 2024 — correctly attributed; the repo file is misnamed "Kim_Suzgun"), ContextEcho, Ross & Lo, Kruthof, Dongre 2025, When Attention Closes, Wu (Meta), CLQT, Kocielnik, Mehri, Hashimoto, Song, Jiang/Peng/Yan, Laban/Always-On survey. None is cited today. Reword the Hong and Liu citations; do not cite ROME for logit attribution.

8. State ContextEcho's compliance-amplifier finding as the null hypothesis for the memory arm and Ross & Lo's heuristic collapse as the reason an opposed-target split is expected; claim novelty of domain and readout, not of injection logic.
9. Degeneracy audit before any cash-shift claim: parse failures (4 in 531,600 steps — publish it), refusals, zero-trade/always-HOLD shares per cell (35.6% of ISFJ-memory flat runs; 47% in crash), and the static-arm cash baseline alongside every contrast (CLQT \$7.11; AI-Trader; KTD-Fin).
10. Cash-mediation table and metric-correlation matrix ($MAS_ISFJ \equiv 1$ – cash; crash MDD effect flips sign controlling cash; ENTJ MAS and bull RG survive the control) — report them before a reviewer does.
11. Portfolio-ratchet decomposition as a finding: per-quartile action rates vs integrated allocation for all personas and arms.
12. P/E leakage: decouple EPS from V or retract "hidden"; report the EPS-divisor sensitivity.
13. Pre-registration of the P2 predictions, stated in the paper; a single metrics pipeline released with the repo (the four-way 0.903/0.942/0.962/0.988 cohort split for one cell must become impossible by construction).
14. Rebase the plan on the current draft and on Cycle 2's error register (34 items) rather than the April draft.

3.4 The plan's "Next Steps", answered

- "Make a research about the previous work" — done: 47 papers read in full (Cycle 2 report §8 and docs/cycle2_reviews/litreview_part1–3.md).
- "Which conference" — COLM main track if the swapped-mandate arm shows content tracking and the paper carries model-level statistics and the positioning above; findings track / strong agents-evaluation workshop otherwise. The decision is made by P2, so run P2 before choosing.
- "The list of experiments we should do as priority" — Shared first steps: P0 truth reconciliation and corrections (1–2 wk, \$0); the cheap four-arm + swapped-mandate package (3–4 wk, ~500–2,000 sims, pre-registered); P1 metric repair + communicated-target track + classifier retrain (2–4 wk). Then Path B: environment rework (time-vs-phase, P/E, δ , RG floors) and the stateful agent design (window/summary/horizon decision) → re-run the grid statefully with the mandate in all arms, keeping the current stateless grid as control (18+ models, months) → M0 on 30–70B open models → M1–M4 as designed, pre-registered with the dissociation outcome explicit (4–8 wk GPU) → controller with behavioural and M2 triggers vs CLQT/Li baselines. Or Path A: P3 = M0 coverage → M1–M3 re-aimed (4–6 wk GPU) → controller → M4 if gated. No compute before P0 lands on either path.

4. One-paragraph summary

Today the paper is a faithful dataset wrapped in an unsupported story: the agents are stateless, the static arm has no mandate, the 4.4× is a half-horizon artefact, and the pooled crash claim fails at model level — but the 17/18-vs-16/18 content effect is exact, model-level significant, replicated, and unowned in its domain. Two revisions are possible. Path B rebuilds the agent (accumulating context, mandate in every arm) and the environment (time decoupled from phase, P/E leak closed) so that Mandate Salience Decay is instantiated for the first time, runs M1–M4 as the plan designed them, keeps the current grid as the stateless control, and is pre-registered to publish the dissociation between accessibility and behaviour that the literature predicts — the bigger paper, months of work, main-track if clean. Path A reframes on the current architecture around content-attribution, with the M-series re-aimed at the directive's position, content and mediation — weeks of work, headline decided by one cheap experiment. Both start with the same corrections, the same cheap causal arms, and the same 16-paper positioning; the current

plan must add the Path-B preconditions, promote the placebo/swap arms from MEDIUM to first, repair rather than derive C_{ideal} , add the architecture rewrite, and be rebased on the current draft.